

MyBERT-local-machine

July 11, 2021

```
[1]: %cd ~/content/  
      %rm -rf ~/content/ActionPrediction4CA  
      %rm -rf ~/content/ActionPredictionBERT  
      !git clone --branch colab_exe https://github.com/jmcraV/ActionPrediction4CA.git
```

```
/home/gian/content  
Clone in 'ActionPrediction4CA' in corso...  
remote: Enumerating objects: 290, done.  
remote: Counting objects: 100% (290/290), done.  
remote: Compressing objects: 100% (206/206), done.  
remote: Total 290 (delta 137), reused 220 (delta 77), pack-reused 0  
Ricezione degli oggetti: 100% (290/290), 42.47 MiB | 10.48 MiB/s, fatto.  
Risoluzione dei delta: 100% (137/137), fatto.
```

```
[2]: %mkdir ~/content/ActionPredictionBERT ActionPredictionBERT/input_data  
      ↳ActionPredictionBERT/extr_output  
      %cd ~/content/ActionPrediction4CA/tools  
      %mv extract_actions_fashion.py ~/content/ActionPredictionBERT  
      %mv action_evaluation.py ~/content/ActionPredictionBERT  
  
      %cd ~/content/ActionPrediction4CA/data/simmc_fashion/  
      %mv fashion_train_dials.json fashion_dev_dials.json  
      ↳fashion_teststd_dials_public.json fashion_metadata.json  
      ↳fashion_devtest_dials.json ~/content/ActionPredictionBERT/input_data  
      %cd ~/content/  
      %rm -rf ./ActionPrediction4CA/
```

```
/home/gian/content/ActionPrediction4CA/tools  
/home/gian/content/ActionPrediction4CA/data/simmc_fashion  
/home/gian/content
```

```
[3]: %cd ~/content/ActionPredictionBERT/  
      !python extract_actions_fashion.py --json_path="./input_data/  
      ↳fashion_train_dials.json ./input_data/fashion_dev_dials.json ./input_data/  
      ↳fashion_devtest_dials.json" --save_root="./extr_output" --metadata_path="./  
      ↳fashion_metadata.json"
```

```
/home/gian/content/ActionPredictionBERT
```

```

Reading: ./input_data/fashion_train_dials.json
Dialogue task Id missing: 3406
Dialogue task Id missing: 3969
Dialogue task Id missing: 4847
Dialogue task Id missing: 321
Dialogue task Id missing: 3455
Dialogue task Id missing: 3414
Saving: ./extr_output/fashion_train_dials_api_calls.json
Reading: ./input_data/fashion_dev_dials.json
Dialogue task Id missing: 2117
Saving: ./extr_output/fashion_dev_dials_api_calls.json
Reading: ./input_data/fashion_devtest_dials.json
Dialogue task Id missing: 9308
Saving: ./extr_output/fashion_devtest_dials_api_calls.json

```

```

[4]: #!/pip install markdown
      #!/pip install transformers
      #!/pip install pandas
      #!/pip install torch

```

```

[5]: #prima parte del fashion_train_dials
      import json
      import pandas as pd
      with open ('./input_data/fashion_train_dials.json','r') as f:
          data= json.load(f)

      result=[]
      row ={}
      for k in data:
          row[k] = data[k]
      # []
      #prima parte del fashion_train_dials
      import json
      import pandas as pd
      with open ('./input_data/fashion_train_dials.json','r') as f:
          data= json.load(f)

      result=[]
      row ={}
      for k in data:
          row[k] = data[k]
      # []

      dialogue_data = pd.json_normalize(row['dialogue_data'])
      type(dialogue_data)
      # dialogue = dialogue_data["dialogue"]
      # for x in dialogue.head(1):

```

```

# display(x)
# #dialogue.head(1)
dialogue_data.head()

dialogue_data = pd.json_normalize(row['dialogue_data'])
type(dialogue_data)
# dialogue = dialogue_data["dialogue"]
# for x in dialogue.head(1):
# display(x)
# #dialogue.head(1)
dialogue_data.head()

```

```

[5]:
                                     dialogue  dialogue_idx  domains \
0  [{'belief_state': [{'act': 'DA:ASK:CHECK:CLOTH...          3094  [fashion]
1  [{'belief_state': [{'act': 'DA:INFORM:PREFER:C...           822  [fashion]
2  [{'belief_state': [{'act': 'DA:REQUEST:GET:CLO...          7411  [fashion]
3  [{'belief_state': [{'act': 'DA:INFORM:DISPREFE...          7029  [fashion]
4  [{'belief_state': [{'act': 'DA:INFORM:DISPREFE...          1506  [fashion]

    dialogue_task_id  dialogue_coref_map.1426  dialogue_coref_map.1429 \
0                1785.0                0.0                1.0
1                1720.0                NaN                NaN
2                2038.0                NaN                NaN
3                2011.0                NaN                NaN
4                1686.0                NaN                NaN

    dialogue_coref_map.708  dialogue_coref_map.712  dialogue_coref_map.2401 \
0                NaN                NaN                NaN
1                0.0                1.0                NaN
2                NaN                NaN                4.0
3                NaN                NaN                NaN
4                NaN                NaN                NaN

    dialogue_coref_map.2402  ...  dialogue_coref_map.2335 \
0                NaN  ...                NaN
1                NaN  ...                NaN
2                0.0  ...                NaN
3                NaN  ...                NaN
4                NaN  ...                NaN

    dialogue_coref_map.713  dialogue_coref_map.1507  dialogue_coref_map.1509 \
0                NaN                NaN                NaN
1                NaN                NaN                NaN
2                NaN                NaN                NaN
3                NaN                NaN                NaN

```

4	NaN	NaN	NaN
	dialogue_coref_map.949	dialogue_coref_map.1137	dialogue_coref_map.1872 \
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN
	dialogue_coref_map.1873	dialogue_coref_map.1753	dialogue_coref_map.834
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

[5 rows x 1648 columns]

```
[6]: #seconda parte del fashion_train_dials
task_mapping = pd.json_normalize(row['task_mapping'])
task_mapping.head()
```

[6]:	task_id	image_ids	focus_image \
0	2042	[2441, 2442, 2443, 2444, 2445, 2446, 2447, 244...	2441
1	2041	[2431, 2432, 2433, 2434, 2435, 2436, 2437, 243...	2431
2	2040	[2421, 2422, 2423, 2424, 2425, 2426, 2427, 242...	2421
3	2039	[2411, 2412, 2413, 2414, 2415, 2416, 2417, 241...	2411
4	2038	[2401, 2402, 2403, 2404, 2405, 2406, 2407, 240...	2401
	memory_images	database_images	
0	[2442, 2443, 2444]	[2445, 2446, 2447, 2448, 2449, 2450]	
1	[2432, 2433, 2434]	[2435, 2436, 2437, 2438, 2439, 2440]	
2	[2422, 2423, 2424]	[2425, 2426, 2427, 2428, 2429, 2430]	
3	[2412, 2413, 2414]	[2415, 2416, 2417, 2418, 2419, 2420]	
4	[2402, 2403, 2404]	[2405, 2406, 2407, 2408, 2409, 2410]	

```
[7]: import pandas as pd
dev_dials_api = pd.read_json('./extr_output/fashion_dev_dials_api_calls.json')
dev_dials_api.head()
```

[7]:	dialog_id	actions \
0	4146	[{'turn_idx': 0, 'action': 'None', 'action_sup...
1	4260	[{'turn_idx': 0, 'action': 'SpecifyInfo', 'act...
2	8022	[{'turn_idx': 0, 'action': 'SearchDatabase', '...
3	4992	[{'turn_idx': 0, 'action': 'None', 'action_sup...
4	5606	[{'turn_idx': 0, 'action': 'None', 'action_sup...

```

                                focus_images
0  [1646, 1646, 1646, 1649, 1649, 1649, 1649]
1                                [2161, 2161, 2161, 2161]
2      [1971, 1972, 1972, 1972, 1977, 1978]
3      [1931, 1931, 1936, 1936, 1936]
4      [1931, 1931, 1931, 1931, 1931]

```

```

[8]: import pandas as pd
devtest_dials_api = pd.read_json('./extr_output/fashion_devtest_dials_api_calls.
    ↪json')
devtest_dials_api.head()

```

```

[8]:   dialog_id                                actions \
0      2494  [{'turn_idx': 0, 'action': 'SearchDatabase', '...'
1      3731  [{'turn_idx': 0, 'action': 'SearchDatabase', '...'
2      8546  [{'turn_idx': 0, 'action': 'SpecifyInfo', 'act...'
3      5590  [{'turn_idx': 0, 'action': 'SearchDatabase', '...'
4      5452  [{'turn_idx': 0, 'action': 'SpecifyInfo', 'act...'

```

```

                                focus_images
0  [1836, 1841, 1841, 1841, 1841]
1  [1676, 1681, 1681, 1683, 1683]
2  [840, 840, 840, 849, 849, 843]
3  [1616, 1618, 1618, 1618, 1618]
4  [2231, 2231, 2231, 2236, 2236]

```

```

[9]: import pandas as pd
import json

def createDataframe(json_file):
    with open(json_file) as f:
        dictftdac = json.load(f)

    data = []

    for e in dictftdac:
        dialog_id = e['dialog_id']
        actions = e['actions']
        focus_images = e['focus_images']

        for a in actions:

            turn_idx = a['turn_idx']
            action = a['action']
            action_supervision = a['action_supervision']
            transcript = a['transcript']
            transcript_annotated = a['transcript_annotated']

```

```

system_transcript = a['system_transcript']
system_transcript_annotated = a['system_transcript_annotated']

row = {
    "dialog_id" : dialog_id,
    'turn_idx' : turn_idx,
    'action' : action,
    'action_supervision' : action_supervision,
    'focus_images' : focus_images,
    'transcript': transcript,
    'transcript_annotated': transcript_annotated,
    'system_transcript': system_transcript,
    'system_transcript_annotated':system_transcript_annotated,
    'previous_transcript': "",
    'previous_system_transcript': ""
}
if (action_supervision != None):
    if 'focus' in action_supervision:
        acsf = {'focus':action_supervision['focus']}
    else:
        acsf = {'focus':None}

    if 'attributes' in action_supervision:
        acsa = {'attributes':action_supervision['attributes']}
    else:
        acsa = {'attributes':[]}
else:
    acsf = {'focus':None}
    acsa = {'attributes':[]}

row.update(acsf)
row.update(acsa)

data.append(row)

# Conservo id turno e risposta sistema per provare a implementare una
→soluzione articolata
df = pd.
→DataFrame(data,columns=['dialog_id','turn_idx','transcript','action','attributes',
→'system_transcript','transcript_annotated','system_transcript_annotated','previous_transcri

return df

```

```

[10]: df_training = createDataframe('./extr_output/fashion_train_dials_api_calls.
    →json')
print("Training: ",len(df_training)," elementi")

```

Training: 21196 elementi

```
[11]: df_validation = createDataframe('./extr_output/fashion_dev_dials_api_calls.  
      ↪json')  
      print("Validation: ",len(df_validation)," elementi")
```

Validation: 3513 elementi

```
[12]: df_test = createDataframe('./extr_output/fashion_devtest_dials_api_calls.json')  
      print("Test: ",len(df_test)," elementi")
```

Test: 5397 elementi

```
[13]: use_next = True
```

```
[14]: #Training  
df_training.sort_values(by=['dialog_id', 'turn_idx'])  
for i in range(1,(len(df_training))):  
    if(i<(len(df_training)) and df_training['dialog_id'][i] ==  
    ↪df_training['dialog_id'][i-1]):  
        df_training.loc[i,'previous_transcript'] = df_training['transcript'][i-1]  
        df_training.loc[i,'previous_system_transcript'] =  
        ↪df_training['system_transcript'][i-1]  
  
#Validation  
df_validation.sort_values(by=['dialog_id', 'turn_idx'])  
for i in range(1,(len(df_validation))):  
    if(i<(len(df_validation)) and df_validation['dialog_id'][i] ==  
    ↪df_validation['dialog_id'][i-1]):  
        df_validation.loc[i,'previous_transcript'] =  
        ↪df_validation['transcript'][i-1]  
        df_validation.loc[i,'previous_system_transcript'] =  
        ↪df_validation['system_transcript'][i-1]  
  
#Evaluation  
df_test.sort_values(by=['dialog_id', 'turn_idx'])  
for i in range(1,(len(df_test))):  
    if(i<(len(df_test)) and df_test['dialog_id'][i] ==  
    ↪df_test['dialog_id'][i-1]):  
        df_test.loc[i,'previous_transcript'] = df_test['transcript'][i-1]  
        df_test.loc[i,'previous_system_transcript'] =  
        ↪df_test['system_transcript'][i-1]
```

```
[15]: from transformers import BertTokenizer  
  
# Load the BERT tokenizer.  
print('Loading BERT tokenizer...')
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased',
    ↪do_lower_case=True)
```

Loading BERT tokenizer...

```
[16]: transcripts_tr = df_training.transcript.values
previous_transcript_tr = df_training.previous_transcript.values
previous_system_transcript_tr = df_training.previous_system_transcript.values
action_labels_tr = df_training.action.values
attributes_labels_tr=df_training.attributes.values

print ("TRAINING DATA:")
# Print the original sentence.
print(' Original: ', transcripts_tr[0])

# Print the sentence split into tokens.
print('Tokenized: ', tokenizer.tokenize(transcripts_tr[0]))

# Print the sentence mapped to token ids.
print('Token IDs: ', tokenizer.convert_tokens_to_ids(tokenizer.
    ↪tokenize(transcripts_tr[0])))
```

TRAINING DATA:

Original: Is there a pattern on this one? It's hard to see in the image.
 Tokenized: ['is', 'there', 'a', 'pattern', 'on', 'this', 'one', '?', 'it', '"',
 's', 'hard', 'to', 'see', 'in', 'the', 'image', '.']
 Token IDs: [2003, 2045, 1037, 5418, 2006, 2023, 2028, 1029, 2009, 1005, 1055,
 2524, 2000, 2156, 1999, 1996, 3746, 1012]

```
[17]: for k in range(0,10):
    print(f"T: {transcripts_tr[k]} | PT: {previous_transcript_tr[k]} | PST:
    ↪{previous_system_transcript_tr[k]}")
```

T: Is there a pattern on this one? It's hard to see in the image. | PT: | PST:
 T: That's fancy. Do you have anything in warmer colors like yellow or red? | PT:
 Is there a pattern on this one? It's hard to see in the image. | PST: I don't
 have any information on the pattern, but it has pointelle embellishments.
 T: Yeah, that sounds good. | PT: That's fancy. Do you have anything in warmer
 colors like yellow or red? | PST: I have a crew neck sweater in red, would you
 like to see it?
 T: Oh, I love that. Please tell me you have a small. | PT: Yeah, that sounds
 good. | PST: This is \$187 from Downtown Stylists with a 3.62 rating.
 T: Yes, please! Thank you for your help with this | PT: Oh, I love that. Please
 tell me you have a small. | PST: It does come in small, shall I put one in your
 cart?
 T: How nice! Does this come in other colors? | PT: | PST:
 T: Oh well. Can you show me a dress that comes in red? | PT: How nice! Does
 this come in other colors? | PST: No, I'm sorry, It comes only in blue.

T: Cute! Do these come in Small? | PT: Oh well. Can you show me a dress that comes in red? | PST: This dress comes in many colors, including a bright red and a pinkish-red. What do you think?

T: Awesome. Would you add a red one in S to my cart please? | PT: Cute! Do these come in Small? | PST: Yes, they do!

T: That's all. Thanks! | PT: Awesome. Would you add a red one in S to my cart please? | PST: The red one is in your cart. Is there anything else I can find for you?

```
[18]: transcripts_vd = df_validation.transcript.values
previous_transcript_vd = df_validation.previous_transcript.values
previous_system_transcript_vd = df_validation.previous_system_transcript.values
action_labels_vd = df_validation.action.values
attributes_labels_vd=df_validation.attributes.values
dialog_ids_vd = df_validation.dialog_id.values
turn_idx_vd = df_validation.turn_idx.values

print ("VALIDATION DATA:")

# Print the original sentence.
print(' Original: ', transcripts_vd[0])

# Print the sentence split into tokens.
print('Tokenized: ', tokenizer.tokenize(transcripts_vd[0]))

# Print the sentence mapped to token ids.
print('Token IDs: ', tokenizer.convert_tokens_to_ids(tokenizer.
→tokenize(transcripts_vd[0])))

# Print the dialog ids.
print(f"Dialog IDs: {dialog_ids_vd[0:20]}")

# Print the turn ids.
print(f"Turn IDs: {turn_idx_vd[0:20]}")
```

VALIDATION DATA:

Original: What's the price of this sweater compared to the other blue and gray one I looked at?

Tokenized: ['what', "'", 's', 'the', 'price', 'of', 'this', 'sweater', 'compared', 'to', 'the', 'other', 'blue', 'and', 'gray', 'one', 'i', 'looked', 'at', '?']

Token IDs: [2054, 1005, 1055, 1996, 3976, 1997, 2023, 14329, 4102, 2000, 1996, 2060, 2630, 1998, 3897, 2028, 1045, 2246, 2012, 1029]

Dialog IDs: [4146 4146 4146 4146 4146 4146 4146 4260 4260 4260 4260 8022 8022 8022

8022 8022 8022 4992 4992 4992]

Turn IDs: [0 1 2 3 4 5 6 0 1 2 3 0 1 2 3 4 5 0 1 2]

```
[19]: transcripts_tst = df_test.transcript.values
previous_transcript_tst = df_test.previous_transcript.values
previous_system_transcript_tst = df_test.previous_system_transcript.values
action_labels_tst = df_test.action.values
attributes_labels_tst= df_test.attributes.values
dialog_ids_tst = df_test.dialog_id.values
turn_idxes_tst = df_test.turn_idx.values

print ("EVALUATION DATA:")

# Print the original sentence.
print(' Original: ', transcripts_tst[0])

# Print the sentence split into tokens.
print('Tokenized: ', tokenizer.tokenize(transcripts_tst[0]))

# Print the sentence mapped to token ids.
print('Token IDs: ', tokenizer.convert_tokens_to_ids(tokenizer.
↳tokenize(transcripts_tst[0])))

# Print the dialog ids.
print(f"Dialog IDs: {dialog_ids_tst[0:20]}")

# Print the turn idxes.
print(f"Turn IDs: {turn_idxes_tst[0:20]}")
```

EVALUATION DATA:

Original: That looks a little too light for what I need, do you have something else with a high customer rating?

Tokenized: ['that', 'looks', 'a', 'little', 'too', 'light', 'for', 'what', 'i', 'need', ',', 'do', 'you', 'have', 'something', 'else', 'with', 'a', 'high', 'customer', 'rating', '?']

Token IDs: [2008, 3504, 1037, 2210, 2205, 2422, 2005, 2054, 1045, 2342, 1010, 2079, 2017, 2031, 2242, 2842, 2007, 1037, 2152, 8013, 5790, 1029]

Dialog IDs: [2494 2494 2494 2494 2494 3731 3731 3731 3731 3731 8546 8546 8546 8546

8546 5590 5590 5590 5590]

Turn IDs: [0 1 2 3 4 0 1 2 3 4 0 1 2 3 4 5 0 1 2 3]

```
[20]: max_len_tr = 0

# For every sentence...
for i in range(0, len(transcripts_tr)):

    # Tokenize the text and add `[CLS]` and `[SEP]` tokens.

    if (previous_transcript_tr[i] != "" and use_next):
```

```

        input_ids = tokenizer.encode(previous_transcript_tr[i]+ " " +_
↪previous_system_transcript_tr[i],transcripts_tr[i], add_special_tokens=True)
    else:
        input_ids = tokenizer.encode(transcripts_tr[i], add_special_tokens=True)

    # Update the maximum sentence length.
    max_len_tr = max(max_len_tr, len(input_ids))

print('Max transcript length for training: ', max_len_tr)

```

Max transcript length for training: 177

```

[21]: max_len_vd = 0

# For every sentence...
for i in range(0,len(transcripts_vd)):

    # Tokenize the text and add `[CLS]` and `[SEP]` tokens.
    if (previous_transcript_vd[i] != "" and use_next):
        input_ids = tokenizer.encode(previous_transcript_vd[i]+ " " +_
↪previous_system_transcript_vd[i],transcripts_vd[i], add_special_tokens=True)
    else:
        input_ids = tokenizer.encode(transcripts_vd[i], add_special_tokens=True)

    # Update the maximum sentence length.
    max_len_vd = max(max_len_vd, len(input_ids))

print('Max transcript length for validation: ', max_len_vd)

```

Max transcript length for validation: 133

```

[22]: max_len_tst = 0

#non sono sicuro che il controllo della lunghezza vada fatto anche sul test_
↪set, dopo la performance non è determinata
#dalla conoscenza del test set?
#è anche vero che in teoria per far funzionare BERT bisogna dargli in pasto dei_
↪dati tokenizzati, quindi in un caso reale il nostro
#model non potrebbe prendere in ingresso del testo non trattato. Nel dubbio ho_
↪controllato le dimensioni

for i in range(0,len(transcripts_tst)):
    # Tokenize the text and add `[CLS]` and `[SEP]` tokens.
    if (previous_transcript_tst[i] != "" and use_next):
        input_ids = tokenizer.encode(previous_transcript_tst[i]+ " " +_
↪previous_system_transcript_tst[i],transcripts_tst[i],_
↪add_special_tokens=True)

```

```

else:
    input_ids = tokenizer.encode(transcripts_tst[i], add_special_tokens=True)

    # Update the maximum sentence length.
    max_len_tst = max(max_len_tst, len(input_ids))

print("Max transcript length for evaluation: ",max_len_tst)

```

Max transcript length for evaluation: 150

```

[23]: max_len = max(max_len_tr, max_len_vd, max_len_tst)

# if (max_len_tr >= max_len_vd):
#     max_len = max_len_tr
# else:
#     max_len = max_len_vd
# if (max_len_tst >= max_len):
#     max_len = max_len_tst

print("La massima lunghezza dei token da gestire è quindi ",max_len)

```

La massima lunghezza dei token da gestire è quindi 177

```

[24]: from sklearn.preprocessing import MultiLabelBinarizer
import numpy as np

mlb = MultiLabelBinarizer()
attributes_labels_all = np.concatenate((attributes_labels_tr,
    ↳ attributes_labels_vd, attributes_labels_tst), axis=None)
attr_yt = mlb.fit_transform(attributes_labels_all)
print(attr_yt[0:15])
print(mlb.inverse_transform(attr_yt[3].reshape(1, -1)))
print(mlb.classes_)
print(f"Totale: {len(attr_yt)}, Training: {len(attributes_labels_tr)},
    ↳ Validation: {len(attributes_labels_vd)}, Evaluation:
    ↳ {len(attributes_labels_tst)}")
attributes_labels_tr_vect = attr_yt[0:len(attributes_labels_tr)]
attributes_labels_vd_vect = attr_yt[len(attributes_labels_tr):
    ↳ (len(attributes_labels_tr)+len(attributes_labels_vd))]
attributes_labels_tst_vect =
    ↳ attr_yt[(len(attributes_labels_tr)+len(attributes_labels_vd)):]
print(f"Training: {len(attributes_labels_tr_vect)}, Validation:
    ↳ {len(attributes_labels_vd_vect)}, Evaluation:
    ↳ {len(attributes_labels_tst_vect)}")

```

```

[[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]

```

```

[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[('availableSizes',)]
['ageRange' 'amountInStock' 'availableSizes' 'brand' 'clothingCategory'
 'clothingStyle' 'color' 'customerRating' 'dressStyle' 'embellishment'
 'forGender' 'forOccasion' 'hasPart' 'hemLength' 'hemStyle' 'info'
 'jacketStyle' 'madeIn' 'material' 'necklineStyle' 'pattern' 'price'
 'sequential' 'size' 'skirtLength' 'skirtStyle' 'sleeveLength'
 'sleeveStyle' 'soldBy' 'sweaterStyle' 'waistStyle' 'warmthRating'
 'waterResistance']
Totale: 30106, Training: 21196, Validation: 3513, Evaluation: 5397
Training: 21196, Validation: 3513, Evaluation: 5397

```

```

[25]: import torch
import tensorflow as tf
torch.backends.cudnn.benchmark = True
torch.backends.cudnn.enabled = True

```

```

[26]: # Tokenize all of the sentences and map the tokens to thier word IDs.

#dobbiamo convertire le nostre labels da string a valori numerici, usiamo il
↳ metodo fornito da sklearn

#TRAINING DATASET
from sklearn import preprocessing

le = preprocessing.LabelEncoder()
action_labels_encoded_tr = le.fit_transform(action_labels_tr)

input_ids_tr = []
attention_masks_tr = []
print(f"{len(df_training)} records to encode.")
# For every sentence...
for i in range(0,len(df_training)):
    # `encode_plus` will:
    #     (1) Tokenize the sentence.

```

```

# (2) Prepend the `[CLS]` token to the start.
# (3) Append the `[SEP]` token to the end.
# (4) Map tokens to their IDs.
# (5) Pad or truncate the sentence to `max_length`
# (6) Create attention masks for [PAD] tokens.

if (previous_transcript_tr[i] != "" and use_next):
    encoded_dict = tokenizer.encode_plus(
        previous_transcript_tr[i] + " " +
        previous_system_transcript_tr[i], # Sentence to encode.
        transcripts_tr[i], #next sentece to encode
        add_special_tokens = True, # Add `[CLS]` and `[SEP]`
        truncation = True,
        max_length = max_len, # Pad & truncate all
        sentences.

        pad_to_max_length = True,
        return_attention_mask = True, # Construct attn. masks.
        return_tensors = 'pt', # Return pytorch tensors.
    )
else:
    encoded_dict = tokenizer.encode_plus(
        transcripts_tr[i], # Sentence to encode.
        add_special_tokens = True, # Add `[CLS]` and `[SEP]`
        truncation = True,
        max_length = max_len, # Pad & truncate all
        sentences.

        pad_to_max_length = True,
        return_attention_mask = True, # Construct attn. masks.
        return_tensors = 'pt', # Return pytorch tensors.
    )

# Add the encoded sentence to the list.
input_ids_tr.append(encoded_dict['input_ids'])

# And its attention mask (simply differentiates padding from non-padding).
attention_masks_tr.append(encoded_dict['attention_mask'])

# Convert the lists into tensors.
input_ids_tr = torch.cat(input_ids_tr, dim=0)
attention_masks_tr = torch.cat(attention_masks_tr, dim=0)
labels_actions_tr = torch.tensor(action_labels_encoded_tr)
labels_attributes_tr = torch.tensor(attributes_labels_tr_vect)

# Print sentence 0, now as a list of IDs.
print ("TRAINING : ")

```



```

input_ids_vd = []
attention_masks_vd = []
print(f"{len(df_validation)} records to encode.")
# For every sentence...
for i in range(0, len(df_validation)):
    # `encode_plus` will:
    #   (1) Tokenize the sentence.
    #   (2) Prepend the `[CLS]` token to the start.
    #   (3) Append the `[SEP]` token to the end.
    #   (4) Map tokens to their IDs.
    #   (5) Pad or truncate the sentence to `max_length`
    #   (6) Create attention masks for [PAD] tokens.

    if (previous_transcript_vd[i] != "" and use_next):
        encoded_dict = tokenizer.encode_plus(
            previous_transcript_vd[i] + " " +
            previous_system_transcript_vd[i], # Sentence to encode.
            transcripts_vd[i], #next sentece to encode
            add_special_tokens = True, # Add '[CLS]' and '[SEP]'
            truncation = True,
            max_length = max_len, # Pad & truncate all
            sentences.

            pad_to_max_length = True,
            return_attention_mask = True, # Construct attn. masks.
            return_tensors = 'pt', # Return pytorch tensors.
        )
    else:
        encoded_dict = tokenizer.encode_plus(
            transcripts_vd[i], # Sentence to encode.
            add_special_tokens = True, # Add '[CLS]' and '[SEP]'
            truncation = True,
            max_length = max_len, # Pad & truncate all
            sentences.

            pad_to_max_length = True,
            return_attention_mask = True, # Construct attn. masks.
            return_tensors = 'pt', # Return pytorch tensors.
        )

    # Add the encoded sentence to the list.
    input_ids_vd.append(encoded_dict['input_ids'])

    # And its attention mask (simply differentiates padding from non-padding).
    attention_masks_vd.append(encoded_dict['attention_mask'])

# Convert the lists into tensors.
input_ids_vd = torch.cat(input_ids_vd, dim=0)
attention_masks_vd = torch.cat(attention_masks_vd, dim=0)

```



```

labels_actions_vd = torch.tensor(action_labels_encoded_vd)
labels_attributes_vd = torch.tensor(attributes_labels_vd_vect)
# Check warning:
# /usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:43: UserWarning:
→To copy construct from a tensor, it is recommended to use sourceTensor.
→clone().detach() or sourceTensor.clone().detach().requires_grad_(True),
→rather than torch.tensor(sourceTensor).
dialog_ids_vd = torch.tensor(dialog_ids_vd)
turn_idxes_vd = torch.tensor(turn_idxes_vd)

# Print sentence 0, now as a list of IDs.
print ("VALIDATION : ")
if (use_next):
    print('Original: ', transcripts_vd[0])
else:
    print('Original: ', transcripts_vd[0])
print('Token IDs:', input_ids_vd[0])
print(f"Dialog IDs: {dialog_ids_vd[0:20]}")
print(f"Turn IDXs: {turn_idxes_vd[0:20]}")

```

3513 records to encode.

VALIDATION :

Original: What's the price of this sweater compared to the other blue and gray one I looked at?

Token IDs: tensor([101, 2054, 1005, 1055, 1996, 3976, 1997, 2023, 14329, 4102,

2000,	1996,	2060,	2630,	1998,	3897,	2028,	1045,	2246,	2012,
1029,	102,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,
0,	0,	0,	0,	0,	0,	0,	0,	0,	0,

Dialog IDs: tensor([4146, 4146, 4146, 4146, 4146, 4146, 4146, 4260, 4260, 4260, 4260, 8022,

8022, 8022, 8022, 8022, 8022, 4992, 4992, 4992])

Turn IDXs: tensor([0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 0, 1, 2, 3, 4, 5, 0, 1, 2])

```

[28]: # Tokenize all of the sentences and map the tokens to thier word IDs.

#dobbiamo convertire le nostre lables da string a valori numerici, usiamo il
↳metodo fornito da sklearn

#VALIDATION DATASET
from sklearn import preprocessing

le = preprocessing.LabelEncoder()
action_labels_encoded_tst = le.fit_transform(action_labels_tst)

input_ids_tst = []
attention_masks_tst = []
print(f"{len(df_test)} records to encode.")
# For every sentence...
for i in range(0, len(df_test)):
# for t in transcripts_tst:
#     `encode_plus` will:
#     (1) Tokenize the sentence.
#     (2) Prepend the `[CLS]` token to the start.
#     (3) Append the `[SEP]` token to the end.
#     (4) Map tokens to their IDs.
#     (5) Pad or truncate the sentence to `max_length`
#     (6) Create attention masks for [PAD] tokens.

#Aggiungere "and False" PER UTILIZZARE sempre la tokenizzazione senza
↳concatenazione
if (previous_transcript_tst[i] != "" and use_next):
    encoded_dict = tokenizer.encode_plus(
        previous_transcript_tst[i] + " " +
↳previous_system_transcript_tst[i], # Sentence to encode.
        transcripts_tst[i], #next sentece to encode
        add_special_tokens = True, # Add '[CLS]' and '[SEP]'
        truncation = True,
        max_length = max_len, # Pad & truncate all
↳sentences.

        pad_to_max_length = True,
        return_attention_mask = True, # Construct attn. masks.
        return_tensors = 'pt', # Return pytorch tensors.
    )
else:
    encoded_dict = tokenizer.encode_plus(
        transcripts_tst[i], # Sentence to encode.
        add_special_tokens = True, # Add '[CLS]' and '[SEP]'
        truncation = True,
        max_length = max_len, # Pad & truncate all
↳sentences.

```



```
[32]: from torch.utils.data import DataLoader, RandomSampler, SequentialSampler

# The DataLoader needs to know our batch size for training, so we specify it
# here. For fine-tuning BERT on a specific task, the authors recommend a batch
# size of 16 or 32.
# With size 32 GeForce RTX 2060 with 6GB run out of memory
batch_size = 12

# Create the DataLoaders for our training and validation sets.
# We'll take training samples in random order.

train_dataloader = DataLoader(
    train_dataset, # The training samples.
    sampler = RandomSampler(train_dataset), # Select batches randomly
    batch_size = batch_size # Trains with this batch size.
)

# For validation the order doesn't matter, so we'll just read them sequentially.
validation_dataloader = DataLoader(
    val_dataset, # The validation samples.
    sampler = SequentialSampler(val_dataset), # Pull out batches
    ↪sequentially.
    batch_size = batch_size # Evaluate with this batch size.
)

#ho controllato nel colab su cui ci basiamo, anche lui usa un Sequential
    ↪Sampler per il dataset di evaluation
evaluation_dataloader = DataLoader(
    tst_dataset, # The validation samples.
    sampler = SequentialSampler(tst_dataset), # Pull out batches
    ↪sequentially.
    batch_size = batch_size # Evaluate with this batch size.
)
```

```
[33]: #DA SISTEMARE
from transformers import BertModel
from torch import nn

class CustomBERTModel(nn.Module):

    def __init__(self):

        super(CustomBERTModel, self).__init__()
        self.bert = BertModel.from_pretrained("bert-base-uncased")
        ### New layers:
        self.linear_intermedio = nn.Linear(768, 256)
```

```

    #provare ad aggiungere ulteriori layer intermedi per ridurre le dimensioni
    ↪fino ad arrivare all'output richiesto
    self.linear_actions = nn.Linear(256, 5)
    self.linear_attributes = nn.Linear(256, len(mlb.classes_)) #num attributi?

def forward(self, ids, mask):
    #controllare che l'output non rappresenti solo lo stato interno dovuto al
    ↪token CLS
    output = self.bert(ids,attention_mask=mask)
    # print(f"Type output{type(output)}")
    # for p in output:
    #     print(p)
    #     print(type(output[p]))
    #     print(output[p])

    #prendiamo il campo last_hidden_state dall'oggetto output; last hidden
    ↪state rappresenta il tensore
    #in uscita dallo step di forward del BertModel
    last_hidden_state_output = output["last_hidden_state"]
    # last_hidden_state has the following shape: (batch_size, sequence_length,
    ↪768)
    #stiamo passando solo il token CLS ai layer successivi
    linear_output_intermedio = self.linear_intermedio(last_hidden_state_output[:
    ↪,0,:].view(-1,768))
    # linear_output_intermedio = self.linear_intermedio(pooled_output)

    linear_output_actions = self.linear_actions(linear_output_intermedio)
    # linear_output_actions = self.sftmx(linear_output_actions)
    # linear_output_actions = nn.functional.softmax(linear_output_actions)
    # Test sigmoid for increasing perplexity performance
    linear_output_actions = torch.sigmoid(linear_output_actions)
    linear_output_attributes = self.linear_attributes(linear_output_intermedio)
    # linear_output_attributes = self.sig(linear_output_attributes)
    linear_output_attributes = torch.sigmoid(linear_output_attributes)

    return {'actions': linear_output_actions, 'attributes':
    ↪linear_output_attributes}

```

```

[34]: #test istanziazione del custom model
model = CustomBERTModel()

# model.bert.config
model.cuda()

```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.predictions.transform.dense.weight', 'cls.seq_relationship.bias', 'cls.predictions.decoder.weight',

```
'cls.predictions.transform.LayerNorm.weight',
'cls.predictions.transform.dense.bias',
'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.bias',
'cls.seq_relationship.weight']
```

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
[34]: CustomBERTModel(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0): BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
          (intermediate): BertIntermediate(
            (dense): Linear(in_features=768, out_features=3072, bias=True)
          )
          (output): BertOutput(
            (dense): Linear(in_features=3072, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
        (1): BertLayer(
          (attention): BertAttention(
```



```

        (self): BertSelfAttention(
          (query): Linear(in_features=768, out_features=768, bias=True)
          (key): Linear(in_features=768, out_features=768, bias=True)
          (value): Linear(in_features=768, out_features=768, bias=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
          (dense): Linear(in_features=768, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
    (intermediate): BertIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
    (output): BertOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
(2): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(3): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(

```

```

        (query): Linear(in_features=768, out_features=768, bias=True)
        (key): Linear(in_features=768, out_features=768, bias=True)
        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
)
(output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
)
)
(4): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
    )
    (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(5): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)

```

```

        (key): Linear(in_features=768, out_features=768, bias=True)
        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
)
(output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
)
)
(6): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
    )
    (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(7): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)

```

```

        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
)
(output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
)
)
(8): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
    )
    (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(9): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)

```

```

        (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
)
(output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
)
)
(10): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
    )
    (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
)
(11): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)

```

```

    )
    (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
)
(output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
)
)
)
)
(pooler): BertPooler(
    (dense): Linear(in_features=768, out_features=768, bias=True)
    (activation): Tanh()
)
)
(linear_intermedio): Linear(in_features=768, out_features=256, bias=True)
(linear_actions): Linear(in_features=256, out_features=5, bias=True)
(linear_attributes): Linear(in_features=256, out_features=33, bias=True)
)

```

```

[35]: # Get all of the model's parameters as a list of tuples.
params = list(model.named_parameters())

print('The BERT model has {:} different named parameters.\n'.
      ↪format(len(params)))

print('==== Embedding Layer ==== \n')

for p in params[0:5]:
    print("{:<55} {:>12}".format(p[0], str(tuple(p[1].size()))))

print('\n==== First Transformer ==== \n')

for p in params[5:21]:
    print("{:<55} {:>12}".format(p[0], str(tuple(p[1].size()))))

print('\n==== Output Layer ==== \n')

for p in params[-4:]:

```

```
print("{:<55} {:>12}".format(p[0], str(tuple(p[1].size()))))
```

The BERT model has 205 different named parameters.

==== Embedding Layer ====

```
bert.embeddings.word_embeddings.weight      (30522, 768)
bert.embeddings.position_embeddings.weight  (512, 768)
bert.embeddings.token_type_embeddings.weight (2, 768)
bert.embeddings.LayerNorm.weight           (768,)
bert.embeddings.LayerNorm.bias             (768,)
```

==== First Transformer ====

```
bert.encoder.layer.0.attention.self.query.weight (768, 768)
bert.encoder.layer.0.attention.self.query.bias  (768,)
bert.encoder.layer.0.attention.self.key.weight  (768, 768)
bert.encoder.layer.0.attention.self.key.bias    (768,)
bert.encoder.layer.0.attention.self.value.weight (768, 768)
bert.encoder.layer.0.attention.self.value.bias  (768,)
bert.encoder.layer.0.attention.output.dense.weight (768, 768)
bert.encoder.layer.0.attention.output.dense.bias (768,)
bert.encoder.layer.0.attention.output.LayerNorm.weight (768,)
bert.encoder.layer.0.attention.output.LayerNorm.bias (768,)
bert.encoder.layer.0.intermediate.dense.weight (3072, 768)
bert.encoder.layer.0.intermediate.dense.bias    (3072,)
bert.encoder.layer.0.output.dense.weight        (768, 3072)
bert.encoder.layer.0.output.dense.bias          (768,)
bert.encoder.layer.0.output.LayerNorm.weight    (768,)
bert.encoder.layer.0.output.LayerNorm.bias      (768,)
```

==== Output Layer ====

```
linear_actions.weight      (5, 256)
linear_actions.bias        (5,)
linear_attributes.weight    (33, 256)
linear_attributes.bias      (33,)
```

```
[36]: from transformers import AdamW
      # Note: AdamW is a class from the huggingface library (as opposed to pytorch)
      # I believe the 'W' stands for 'Weight Decay fix'
      optimizer = AdamW(model.parameters(),
                        lr = 5e-5, # args.learning_rate - default is 5e-5
                        eps = 1e-8 # args.adam_epsilon - default is 1e-8.
                        )
```

```
[37]: from transformers import get_linear_schedule_with_warmup

# Number of training epochs. The BERT authors recommend between 2 and 4.
# We chose to run for 4, but we'll see later that this may be over-fitting the
# training data.
epochs = 4

# Total number of training steps is [number of batches] x [number of epochs].
# (Note that this is not the same as the number of training samples).
total_steps = len(train_dataloader) * epochs

# Create the learning rate scheduler.
scheduler = get_linear_schedule_with_warmup(optimizer,
                                             num_warmup_steps = 0, # Default
                                             ↪value in run_glue.py
                                             num_training_steps = total_steps)
```

```
[38]: import numpy as np

# Function to calculate the accuracy of our predictions vs labels
def flat_accuracy_actions(preds, labels):
    #print(f"[FA] preds: {preds} / labels: {labels}")
    #print(f"[FA-Actions] {type(preds)} {type(labels)}")
    pred_flat = np.argmax(preds, axis=1).flatten()
    labels_flat = labels.flatten()
    return {'matched': np.sum(pred_flat == labels_flat), 'counts': ↪
    ↪len(labels_flat)}

def flat_accuracy_attributes(preds, labels):
    #print(f"[FA-Attributes] {type(preds)} {type(labels)}")
    tot_preds = preds.shape[0]
    preds_int = np rint(preds)
    tot_eq = 0
    for i in range(tot_preds):
        comparison = preds_int[i] == labels[i]
        if comparison.all():
            tot_eq += 1
    return {'matched': tot_eq, 'counts' : tot_preds}
```

```
[39]: import time
import datetime

def format_time(elapsed):
    '''
    Takes a time in seconds and returns a string hh:mm:ss
    '''
    # Round to the nearest second.
```



```

elapsed_rounded = int(round((elapsed)))

# Format as hh:mm:ss
return str(datetime.timedelta(seconds=elapsed_rounded))

```

```

[40]: from torch import nn
      # Loss function definition
      def MyBERT_loss(logits, actions_labels, attributes_labels):
          actions_logits = logits['actions']
          attributes_logits = logits['attributes']
          loss_actions_fn = nn.CrossEntropyLoss()
          loss_attributes_fn = nn.BCELoss()
          loss_actions = loss_actions_fn(actions_logits, actions_labels)
          loss_attributes = loss_attributes_fn(attributes_logits, attributes_labels.
      ↪float())
          return loss_actions + loss_attributes

```

```

[41]: #from GPUUtil import showUtilization as gpu_usage

```

```

[42]: import random
      import numpy as np
      import action_evaluation as evaluation
      import json
      from GPUUtil import showUtilization as gpu_usage
      with open('./extr_output/fashion_dev_dials_api_calls.json') as f:
          dev_dials = json.load(f)

      # This training code is based on the `run_glue.py` script here:
      # https://github.com/huggingface/transformers/blob/
      ↪5bfcd0485ece086ebcbed2d008813037968a9e58/examples/run_glue.py#L128

      # Set the seed value all over the place to make this reproducible.
      seed_val = 24

      random.seed(seed_val)
      np.random.seed(seed_val)
      torch.manual_seed(seed_val)
      torch.cuda.manual_seed_all(seed_val)

      # We'll store a number of quantities such as training and validation loss,
      # validation accuracy, and timings.
      training_stats = []

      # Measure the total training time for the whole run.
      total_t0 = time.time()

      test_batch = []

```

```

# For each epoch...
for epoch_i in range(0, epochs):

    # =====
    #           Training
    # =====

    # Perform one full pass over the training set.

    print("")
    print('=====Epoch {:} / {:} ====='.format(epoch_i + 1, epochs))
    print('Training...')

    # Measure how long the training epoch takes.
    t0 = time.time()

    # Reset the total loss for this epoch.
    total_train_loss = 0

    # Put the model into training mode. Don't be misled--the call to
    # `train` just changes the *mode*, it doesn't *perform* the training.
    # `dropout` and `batchnorm` layers behave differently during training
    # vs. test (source: https://stackoverflow.com/questions/51433378/
    ↪what-does-model-train-do-in-pytorch)
    print("GPU before train")
    gpu_usage()

    model.train()
    print("GPU after train")
    gpu_usage()

    # For each batch of training data...
    for step, batch in enumerate(train_dataloader):

        # Progress update every 40 batches.
        if step % 40 == 0 and not step == 0:
            # Calculate elapsed time in minutes.
            elapsed = format_time(time.time() - t0)

            # Report progress.
            print('  Batch {:>5,} of {:>5,}.    Elapsed: {:}'.format(step,
            ↪len(train_dataloader), elapsed))
            #DEBUG -- da levare
            #break
            if step == 40 or step % 400 == 0:
                gpu_usage()

```

```

# Unpack this training batch from our dataloader.
#
# As we unpack the batch, we'll also copy each tensor to the GPU using
→ the
# `to` method.
#
# `batch` contains three pytorch tensors:
# [0]: input ids
# [1]: attention masks
# [2]: actions labels
# [3]: attributes labels
b_input_ids = batch[0].to(device)
b_input_mask = batch[1].to(device)
b_labels_actions = batch[2].to(device)
b_labels_attributes = batch[3].to(device)

# Always clear any previously calculated gradients before performing a
# backward pass. PyTorch doesn't do this automatically because
# accumulating the gradients is "convenient while training RNNs".
# (source: https://stackoverflow.com/questions/48001598/
→ why-do-we-need-to-call-zero-grad-in-pytorch)
model.zero_grad()

# Perform a forward pass (evaluate the model on this training batch).
# In PyTorch, calling `model` will in turn call the model's `forward`
# function and pass down the arguments. The `forward` function is
# documented here:
# https://huggingface.co/transformers/model\_doc/bert.
→ html#bertforsequenceclassification
# The results are returned in a results object, documented here:
# https://huggingface.co/transformers/main\_classes/output.
→ html#transformers.modeling\_outputs.SequenceClassifierOutput
# Specifically, we'll get the loss (because we provided labels) and the
# "logits"--the model outputs prior to activation.
result = model(b_input_ids,
               mask=b_input_mask)

loss = MyBERT_loss(result, b_labels_actions, b_labels_attributes)

# Accumulate the training loss over all of the batches so that we can
# calculate the average loss at the end. `loss` is a Tensor containing a
# single value; the `.item()` function just returns the Python value
# from the tensor.
total_train_loss += loss.item()

# Perform a backward pass to calculate the gradients.

```

```

    loss.backward()

    # Clip the norm of the gradients to 1.0.
    # This is to help prevent the "exploding gradients" problem.
    → transformers import BertModel, BertConfig
    torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)

    # Update parameters and take a step using the computed gradient.
    # The optimizer dictates the "update rule"--how the parameters are
    # modified based on their gradients, the learning rate, etc.
    optimizer.step()

    # Update the learning rate.
    scheduler.step()

    print(f"End of epoch {epoch_i}")
    gpu_usage()

    # Calculate the average loss over all of the batches.
    avg_train_loss = total_train_loss / len(train_dataloader)

    # Measure how long this epoch took.
    training_time = format_time(time.time() - t0)

    print("")
    print("  Average training loss: {0:.2f}".format(avg_train_loss))
    print("  Training epoch took: {:}".format(training_time))

    # =====
    #                      Validation
    # =====
    # After the completion of each training epoch, measure our performance on
    # our validation set.

    print("")
    print("Running Validation...")

    t0 = time.time()

    # Put the model in evaluation mode--the dropout layers behave differently
    # during evaluation.
    model.eval()

    # Tracking variables
    total_eval_accuracy_classification = { 'matched': 0, 'counts': 0 }
    total_eval_accuracy_multilabel = { 'matched': 0, 'counts': 0 }
    total_eval_loss = 0

```

```

nb_eval_steps = 0

batch_number = 0

# Dictionary for action_evaluation
model_actions = {}

# Evaluate data for one epoch
for batch in validation_dataloader:

    batch_number += 1

    # Unpack this training batch from our dataloader.
    #
    # As we unpack the batch, we'll also copy each tensor to the GPU using
    # the `to` method.
    #
    # `batch` contains three pytorch tensors:
    # [0]: input ids
    # [1]: attention masks
    # [2]: labels
    b_input_ids = batch[0].to(device)
    b_input_mask = batch[1].to(device)
    b_labels_actions = batch[2].to(device)
    b_labels_attributes = batch[3].to(device)
    b_dialog_ids = batch[4].to(device).detach().cpu().numpy()
    b_turn_idx = batch[5].to(device).detach().cpu().numpy()

    # Tell pytorch not to bother with constructing the compute graph during
    # the forward pass, since this is only needed for backprop (training).
    with torch.no_grad():

        # Forward pass, calculate logit predictions.
        # token_type_ids is the same as the "segment ids", which
        # differentiates sentence 1 and 2 in 2-sentence tasks.
        result = model(b_input_ids,
                        mask=b_input_mask)

        # Get the loss and "logits" output by the model. The "logits" are the
        # output values prior to applying an activation function like the
        # softmax.
        loss = MyBERT_loss(result, b_labels_actions, b_labels_attributes)

        # Accumulate the validation loss.
        total_eval_loss += loss.item()

    # Move logits and labels to CPU

```

```

# logits = logits.detach().cpu().numpy()
# label_ids = b_labels.to('cpu').numpy()

actions_logits_foracc=result['actions'].detach().cpu().numpy()
attributes_logits_foracc=result['attributes'].detach().cpu().numpy()
actions_labels_foracc= b_labels_actions.to('cpu').numpy()
attributes_labels_foracc =b_labels_attributes.to('cpu').numpy()

#TODO: definire la nostra funzione di accuracy

# Calculate the accuracy for this batch of test sentences, and
# accumulate it over all batches.
accuracy_classification = flat_accuracy_actions(actions_logits_foracc,
↪actions_labels_foracc)
accuracy_multilabel =
↪flat_accuracy_attributes(attributes_logits_foracc, attributes_labels_foracc)

total_eval_accuracy_classification['matched'] +=
↪accuracy_classification['matched']
total_eval_accuracy_classification['counts'] +=
↪accuracy_classification['counts']
total_eval_accuracy_multilabel['matched'] +=
↪accuracy_multilabel['matched']
total_eval_accuracy_multilabel['counts'] +=
↪accuracy_multilabel['counts']
# Salvo dati elaborazione batch per debug/analisi
test_batch.append({
    'epoch' : epoch_i + 1,
    'batchnum' : batch_number,
    'actions_logits' : actions_logits_foracc,
    'actions_labels' : actions_labels_foracc,
    'attributes_logits' : attributes_logits_foracc,
    'attributes_labels' : attributes_labels_foracc,
    'accuracy_classification' : accuracy_classification,
    'accuracy_multilabel' : accuracy_multilabel,
})

# Fill dictionary for action_evaluation
for el_i in range(len(actions_logits_foracc)):
    dialog_id = b_dialog_ids[el_i]
    action_log_prob = {}
    for act_i in range(len(actions_logits_foracc[el_i])):
        #todo: controllare che la probabilità predetta sia in scala
↪logaritmica (?? potrebbe essere fonte di errori)

```

```

        action_log_prob[le.classes_[act_i]] = np.
→log(actions_logits_foracc[el_i][act_i])
        #attributes = {}
        attributes = []
        #attributes_list = np.rint(attributes_logits_foracc[el_i])
        attributes_list = np.array(attributes_logits_foracc[el_i])
        for attr in range(len(attributes_list)):
            attribute = mlb.classes_[attr]
            #attributes[mlb.classes_[attr]] = attributes_list[attr]
            if attributes_list[attr] >= 0.5:
                attributes.append(attribute)
        prediction = {
            'action': le.classes_[np.argmax(actions_logits_foracc[el_i])],
            'action_log_prob': action_log_prob,
            'attributes': {'attributes': attributes},
            'turn_id': b_turn_idx[el_i]
        }
        if dialog_id in model_actions:
            model_actions[dialog_id]['predictions'].append(prediction)
        else:
            predictions = list()
            predictions.append(prediction)
            model_actions[dialog_id] = {
                'dialog_id': dialog_id,
                'predictions': predictions
            }

    # Report the final accuracy for this validation

    #avg_val_accuracy_classification = total_eval_accuracy_classification /
→len(validation_data_loader)
    #avg_val_accuracy_multilabel = total_eval_accuracy_multilabel /
→len(validation_data_loader)
    avg_val_accuracy_classification =
→total_eval_accuracy_classification['matched'] /
→total_eval_accuracy_classification['counts']
    avg_val_accuracy_multilabel = total_eval_accuracy_multilabel['matched'] /
→total_eval_accuracy_multilabel['counts']
    print(" Accuracy for classification (actions): {0:.4f}".
→format(avg_val_accuracy_classification))
    print(" Accuracy for multilabel-classification (attributes): {0:.4f}".
→format(avg_val_accuracy_multilabel))

    # Reference implementation: evaluation of action prediction along with
→attributes

```

```

    metrics = evaluation.evaluate_action_prediction(dev_dials, model_actions.
↪values())
    # print("model_actions passed to the evaluator:")
    # for v in model_actions.values():
    #     print(v)
    print("*****")
    print("Reference evaluation metrics:")
    print(metrics)

    # Calculate the average loss over all of the batches.
    avg_val_loss = total_eval_loss / len(validation_data_loader)

    # Measure how long the validation run took.
    validation_time = format_time(time.time() - t0)

    print("  Validation Loss: {0:.4f}".format(avg_val_loss))
    print("  Validation took: {:}".format(validation_time))

    # Record all statistics from this epoch.
    training_stats.append(
        {
            'epoch': epoch_i + 1,
            'Training Loss': avg_train_loss,
            'Valid. Loss': avg_val_loss,
            'Valid. Accur. class.': avg_val_accuracy_classification,
            'Valid. Accur. mult.label': avg_val_accuracy_multilabel,
            'Training Time': training_time,
            'Validation Time': validation_time
        }
    )

    print("")
    print("Training complete!")

    print("Total training took {:} (h:mm:ss)".format(format_time(time.
↪time()-total_t0)))

```

===== Epoch 1 / 4 =====

Training...

GPU before train

ID	GPU	MEM
0	4%	29%

0	4%	29%
---	----	-----

GPU after train

ID	GPU	MEM
----	-----	-----

0 4% 29%	
Batch 40 of 1,767.	Elapsed: 0:00:12.
ID GPU MEM	

0 99% 88%	
Batch 80 of 1,767.	Elapsed: 0:00:23.
Batch 120 of 1,767.	Elapsed: 0:00:34.
Batch 160 of 1,767.	Elapsed: 0:00:46.
Batch 200 of 1,767.	Elapsed: 0:00:57.
Batch 240 of 1,767.	Elapsed: 0:01:08.
Batch 280 of 1,767.	Elapsed: 0:01:20.
Batch 320 of 1,767.	Elapsed: 0:01:31.
Batch 360 of 1,767.	Elapsed: 0:01:42.
Batch 400 of 1,767.	Elapsed: 0:01:54.
ID GPU MEM	

0 99% 87%	
Batch 440 of 1,767.	Elapsed: 0:02:05.
Batch 480 of 1,767.	Elapsed: 0:02:17.
Batch 520 of 1,767.	Elapsed: 0:02:29.
Batch 560 of 1,767.	Elapsed: 0:02:41.
Batch 600 of 1,767.	Elapsed: 0:02:53.
Batch 640 of 1,767.	Elapsed: 0:03:05.
Batch 680 of 1,767.	Elapsed: 0:03:16.
Batch 720 of 1,767.	Elapsed: 0:03:28.
Batch 760 of 1,767.	Elapsed: 0:03:40.
Batch 800 of 1,767.	Elapsed: 0:03:52.
ID GPU MEM	

0 99% 89%	
Batch 840 of 1,767.	Elapsed: 0:04:04.
Batch 880 of 1,767.	Elapsed: 0:04:16.
Batch 920 of 1,767.	Elapsed: 0:04:28.
Batch 960 of 1,767.	Elapsed: 0:04:40.
Batch 1,000 of 1,767.	Elapsed: 0:04:52.
Batch 1,040 of 1,767.	Elapsed: 0:05:04.
Batch 1,080 of 1,767.	Elapsed: 0:05:16.
Batch 1,120 of 1,767.	Elapsed: 0:05:28.
Batch 1,160 of 1,767.	Elapsed: 0:05:40.
Batch 1,200 of 1,767.	Elapsed: 0:05:52.
ID GPU MEM	

0 99% 90%	
Batch 1,240 of 1,767.	Elapsed: 0:06:04.
Batch 1,280 of 1,767.	Elapsed: 0:06:15.
Batch 1,320 of 1,767.	Elapsed: 0:06:27.
Batch 1,360 of 1,767.	Elapsed: 0:06:38.
Batch 1,400 of 1,767.	Elapsed: 0:06:50.

```

Batch 1,440 of 1,767. Elapsed: 0:07:02.
Batch 1,480 of 1,767. Elapsed: 0:07:13.
Batch 1,520 of 1,767. Elapsed: 0:07:25.
Batch 1,560 of 1,767. Elapsed: 0:07:37.
Batch 1,600 of 1,767. Elapsed: 0:07:48.
| ID | GPU | MEM |
-----
| 0 | 99% | 88% |
Batch 1,640 of 1,767. Elapsed: 0:08:00.
Batch 1,680 of 1,767. Elapsed: 0:08:13.
Batch 1,720 of 1,767. Elapsed: 0:08:24.
Batch 1,760 of 1,767. Elapsed: 0:08:36.
End of epoch 0
| ID | GPU | MEM |
-----
| 0 | 99% | 90% |

Average training loss: 1.13
Training epoch took: 0:08:38

Running Validation...
Accuracy for classification (actions): 0.8451
Accuracy for multilabel-classification (attributes): 0.8785
#Instances evaluated API: 3513
*****
Reference evaluation metrics:
{'action_accuracy': 0.8451465983489894, 'action_perplexity': 2.061376839726222,
'attribute_accuracy': 0.6787099391419993, 'confusion_matrix': array([[4.720e+02,
4.600e+01, 7.000e+00, 1.000e+00, 1.600e+01],
[1.000e+01, 6.530e+02, 2.600e+01, 4.000e+00, 9.000e+00],
[1.000e+01, 1.500e+02, 5.090e+02, 4.700e+01, 1.500e+01],
[0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00],
[0.000e+00, 5.900e+01, 8.200e+01, 6.200e+01, 1.335e+03]])}
Validation Loss: 1.0664
Validation took: 0:00:28

===== Epoch 2 / 4 =====
Training...
GPU before train
| ID | GPU | MEM |
-----
| 0 | 98% | 89% |
GPU after train
| ID | GPU | MEM |
-----
| 0 | 98% | 89% |
Batch 40 of 1,767. Elapsed: 0:00:12.
| ID | GPU | MEM |

```

```

-----
| 0 | 100% | 90% |
Batch 80 of 1,767. Elapsed: 0:00:24.
Batch 120 of 1,767. Elapsed: 0:00:35.
Batch 160 of 1,767. Elapsed: 0:00:47.
Batch 200 of 1,767. Elapsed: 0:00:59.
Batch 240 of 1,767. Elapsed: 0:01:11.
Batch 280 of 1,767. Elapsed: 0:01:23.
Batch 320 of 1,767. Elapsed: 0:01:35.
Batch 360 of 1,767. Elapsed: 0:01:47.
Batch 400 of 1,767. Elapsed: 0:01:59.

```

```

| ID | GPU | MEM |
-----

```

```

| 0 | 99% | 91% |
Batch 440 of 1,767. Elapsed: 0:02:11.
Batch 480 of 1,767. Elapsed: 0:02:22.
Batch 520 of 1,767. Elapsed: 0:02:34.
Batch 560 of 1,767. Elapsed: 0:02:46.
Batch 600 of 1,767. Elapsed: 0:02:58.
Batch 640 of 1,767. Elapsed: 0:03:09.
Batch 680 of 1,767. Elapsed: 0:03:21.
Batch 720 of 1,767. Elapsed: 0:03:32.
Batch 760 of 1,767. Elapsed: 0:03:44.
Batch 800 of 1,767. Elapsed: 0:03:55.

```

```

| ID | GPU | MEM |
-----

```

```

| 0 | 95% | 89% |
Batch 840 of 1,767. Elapsed: 0:04:07.
Batch 880 of 1,767. Elapsed: 0:04:18.
Batch 920 of 1,767. Elapsed: 0:04:30.
Batch 960 of 1,767. Elapsed: 0:04:42.
Batch 1,000 of 1,767. Elapsed: 0:04:53.
Batch 1,040 of 1,767. Elapsed: 0:05:05.
Batch 1,080 of 1,767. Elapsed: 0:05:16.
Batch 1,120 of 1,767. Elapsed: 0:05:28.
Batch 1,160 of 1,767. Elapsed: 0:05:39.
Batch 1,200 of 1,767. Elapsed: 0:05:51.

```

```

| ID | GPU | MEM |
-----

```

```

| 0 | 99% | 89% |
Batch 1,240 of 1,767. Elapsed: 0:06:03.
Batch 1,280 of 1,767. Elapsed: 0:06:15.
Batch 1,320 of 1,767. Elapsed: 0:06:26.
Batch 1,360 of 1,767. Elapsed: 0:06:38.
Batch 1,400 of 1,767. Elapsed: 0:06:49.
Batch 1,440 of 1,767. Elapsed: 0:07:02.
Batch 1,480 of 1,767. Elapsed: 0:07:14.
Batch 1,520 of 1,767. Elapsed: 0:07:26.

```

```

Batch 1,560 of 1,767. Elapsed: 0:07:37.
Batch 1,600 of 1,767. Elapsed: 0:07:49.
| ID | GPU | MEM |

```

```

-----
| 0 | 100% | 90% |
Batch 1,640 of 1,767. Elapsed: 0:08:01.
Batch 1,680 of 1,767. Elapsed: 0:08:12.
Batch 1,720 of 1,767. Elapsed: 0:08:25.
Batch 1,760 of 1,767. Elapsed: 0:08:36.

```

End of epoch 1

```

| ID | GPU | MEM |

```

```

-----
| 0 | 99% | 89% |

```

Average training loss: 1.06

Training epoch took: 0:08:38

Running Validation...

Accuracy for classification (actions): 0.8488

Accuracy for multilabel-classification (attributes): 0.9209

#Instances evaluated API: 3513

Reference evaluation metrics:

```

{'action_accuracy': 0.8488471391972673, 'action_perplexity': 2.3182883517974853,
'attribute_accuracy': 0.7356263777185521, 'confusion_matrix': array([[4.680e+02,
3.600e+01, 4.000e+00, 0.000e+00, 7.000e+00],
[1.400e+01, 6.790e+02, 3.700e+01, 1.000e+01, 1.000e+01],
[9.000e+00, 1.370e+02, 4.840e+02, 4.100e+01, 7.000e+00],
[0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00],
[1.000e+00, 5.600e+01, 9.900e+01, 6.300e+01, 1.351e+03]])}

```

Validation Loss: 1.0584

Validation took: 0:00:27

===== Epoch 3 / 4 =====

Training...

GPU before train

```

| ID | GPU | MEM |

```

```

-----
| 0 | 97% | 88% |

```

GPU after train

```

| ID | GPU | MEM |

```

```

-----
| 0 | 97% | 88% |

```

```

Batch 40 of 1,767. Elapsed: 0:00:12.

```

```

| ID | GPU | MEM |

```

```

-----
| 0 | 99% | 88% |

```

```

Batch 80 of 1,767. Elapsed: 0:00:24.

```

Batch	120	of	1,767.	Elapsed:	0:00:35.
Batch	160	of	1,767.	Elapsed:	0:00:47.
Batch	200	of	1,767.	Elapsed:	0:00:59.
Batch	240	of	1,767.	Elapsed:	0:01:10.
Batch	280	of	1,767.	Elapsed:	0:01:22.
Batch	320	of	1,767.	Elapsed:	0:01:33.
Batch	360	of	1,767.	Elapsed:	0:01:45.
Batch	400	of	1,767.	Elapsed:	0:01:56.

ID GPU MEM

0 99% 88%

Batch	440	of	1,767.	Elapsed:	0:02:08.
Batch	480	of	1,767.	Elapsed:	0:02:20.
Batch	520	of	1,767.	Elapsed:	0:02:31.
Batch	560	of	1,767.	Elapsed:	0:02:43.
Batch	600	of	1,767.	Elapsed:	0:02:54.
Batch	640	of	1,767.	Elapsed:	0:03:06.
Batch	680	of	1,767.	Elapsed:	0:03:17.
Batch	720	of	1,767.	Elapsed:	0:03:29.
Batch	760	of	1,767.	Elapsed:	0:03:40.
Batch	800	of	1,767.	Elapsed:	0:03:52.

ID GPU MEM

0 100% 89%

Batch	840	of	1,767.	Elapsed:	0:04:04.
Batch	880	of	1,767.	Elapsed:	0:04:16.
Batch	920	of	1,767.	Elapsed:	0:04:28.
Batch	960	of	1,767.	Elapsed:	0:04:40.
Batch	1,000	of	1,767.	Elapsed:	0:04:52.
Batch	1,040	of	1,767.	Elapsed:	0:05:04.
Batch	1,080	of	1,767.	Elapsed:	0:05:15.
Batch	1,120	of	1,767.	Elapsed:	0:05:27.
Batch	1,160	of	1,767.	Elapsed:	0:05:38.
Batch	1,200	of	1,767.	Elapsed:	0:05:50.

ID GPU MEM

0 99% 89%

Batch	1,240	of	1,767.	Elapsed:	0:06:02.
Batch	1,280	of	1,767.	Elapsed:	0:06:14.
Batch	1,320	of	1,767.	Elapsed:	0:06:25.
Batch	1,360	of	1,767.	Elapsed:	0:06:37.
Batch	1,400	of	1,767.	Elapsed:	0:06:49.
Batch	1,440	of	1,767.	Elapsed:	0:07:00.
Batch	1,480	of	1,767.	Elapsed:	0:07:12.
Batch	1,520	of	1,767.	Elapsed:	0:07:24.
Batch	1,560	of	1,767.	Elapsed:	0:07:36.
Batch	1,600	of	1,767.	Elapsed:	0:07:48.

ID GPU MEM

```

-----
| 0 | 99% | 89% |
Batch 1,640 of 1,767. Elapsed: 0:07:59.
Batch 1,680 of 1,767. Elapsed: 0:08:11.
Batch 1,720 of 1,767. Elapsed: 0:08:23.
Batch 1,760 of 1,767. Elapsed: 0:08:35.
End of epoch 2
| ID | GPU | MEM |
-----
| 0 | 99% | 89% |

Average training loss: 1.04
Training epoch took: 0:08:37

Running Validation...
Accuracy for classification (actions): 0.8514
Accuracy for multilabel-classification (attributes): 0.9243
#Instances evaluated API: 3513
*****
Reference evaluation metrics:
{'action_accuracy': 0.8514090520922288, 'action_perplexity': 2.4830429491963275,
'attribute_accuracy': 0.7453504514504714, 'confusion_matrix': array([[ 468.,
36.,    6.,    0.,    4.],
[ 10., 695.,  70.,  14.,  16.],
[  8., 125., 483.,  45.,  10.],
[  0.,  0.,  0.,  0.,  0.],
[  6.,  52.,  65.,  55., 1345.]])}
Validation Loss: 1.0541
Validation took: 0:00:27

===== Epoch 4 / 4 =====
Training...
GPU before train
| ID | GPU | MEM |
-----
| 0 | 98% | 89% |
GPU after train
| ID | GPU | MEM |
-----
| 0 | 98% | 89% |
Batch 40 of 1,767. Elapsed: 0:00:12.
| ID | GPU | MEM |
-----
| 0 | 99% | 89% |
Batch 80 of 1,767. Elapsed: 0:00:24.
Batch 120 of 1,767. Elapsed: 0:00:36.
Batch 160 of 1,767. Elapsed: 0:00:47.
Batch 200 of 1,767. Elapsed: 0:00:59.

```

Batch	240	of	1,767.	Elapsed:	0:01:11.
Batch	280	of	1,767.	Elapsed:	0:01:23.
Batch	320	of	1,767.	Elapsed:	0:01:35.
Batch	360	of	1,767.	Elapsed:	0:01:46.
Batch	400	of	1,767.	Elapsed:	0:01:58.
ID	GPU	MEM			

0	99%	89%			
Batch	440	of	1,767.	Elapsed:	0:02:10.
Batch	480	of	1,767.	Elapsed:	0:02:22.
Batch	520	of	1,767.	Elapsed:	0:02:34.
Batch	560	of	1,767.	Elapsed:	0:02:46.
Batch	600	of	1,767.	Elapsed:	0:02:57.
Batch	640	of	1,767.	Elapsed:	0:03:09.
Batch	680	of	1,767.	Elapsed:	0:03:21.
Batch	720	of	1,767.	Elapsed:	0:03:33.
Batch	760	of	1,767.	Elapsed:	0:03:44.
Batch	800	of	1,767.	Elapsed:	0:03:56.
ID	GPU	MEM			

0	99%	89%			
Batch	840	of	1,767.	Elapsed:	0:04:08.
Batch	880	of	1,767.	Elapsed:	0:04:20.
Batch	920	of	1,767.	Elapsed:	0:04:31.
Batch	960	of	1,767.	Elapsed:	0:04:43.
Batch	1,000	of	1,767.	Elapsed:	0:04:55.
Batch	1,040	of	1,767.	Elapsed:	0:05:06.
Batch	1,080	of	1,767.	Elapsed:	0:05:18.
Batch	1,120	of	1,767.	Elapsed:	0:05:29.
Batch	1,160	of	1,767.	Elapsed:	0:05:41.
Batch	1,200	of	1,767.	Elapsed:	0:05:53.
ID	GPU	MEM			

0	99%	89%			
Batch	1,240	of	1,767.	Elapsed:	0:06:04.
Batch	1,280	of	1,767.	Elapsed:	0:06:16.
Batch	1,320	of	1,767.	Elapsed:	0:06:28.
Batch	1,360	of	1,767.	Elapsed:	0:06:39.
Batch	1,400	of	1,767.	Elapsed:	0:06:51.
Batch	1,440	of	1,767.	Elapsed:	0:07:03.
Batch	1,480	of	1,767.	Elapsed:	0:07:15.
Batch	1,520	of	1,767.	Elapsed:	0:07:26.
Batch	1,560	of	1,767.	Elapsed:	0:07:38.
Batch	1,600	of	1,767.	Elapsed:	0:07:50.
ID	GPU	MEM			

0	99%	89%			
Batch	1,640	of	1,767.	Elapsed:	0:08:02.

```

Batch 1,680 of 1,767. Elapsed: 0:08:14.
Batch 1,720 of 1,767. Elapsed: 0:08:25.
Batch 1,760 of 1,767. Elapsed: 0:08:37.
End of epoch 3
| ID | GPU | MEM |
-----
| 0 | 99% | 89% |

Average training loss: 1.03
Training epoch took: 0:08:39

Running Validation...
Accuracy for classification (actions): 0.8565
Accuracy for multilabel-classification (attributes): 0.9280
#Instances evaluated API: 3513
*****
Reference evaluation metrics:
{'action_accuracy': 0.856532877882152, 'action_perplexity': 2.810755603660233,
'attribute_accuracy': 0.7616486252087795, 'confusion_matrix': array([[ 469.,
33.,   4.,   0.,   4.],
[  8., 696.,  54.,  10.,  15.],
[ 10., 133., 502.,  48.,  14.],
[  0.,   0.,   0.,   0.,   0.],
[  5.,  46.,  64.,  56., 1342.]])}
Validation Loss: 1.0542
Validation took: 0:00:27

Training complete!
Total training took 0:36:21 (h:mm:ss)

```

```

[43]: #Prediction on test set
#quale modello gli viene passato? da controllare se BERT da solo riesce a
↳tenere traccia del modello che ha dato l'epoca migliore

with open('./extr_output/fashion_devtest_dials_api_calls.json') as f:
    devtest_dials = json.load(f)

# Tracking variables
total_eval_accuracy_classification = { 'matched': 0, 'counts': 0}
total_eval_accuracy_multilabel = { 'matched': 0, 'counts': 0}

model_actions = {}
# Put model in evaluation mode
model.eval()

for batch in evaluation_dataloader:

```



```

# Unpack this training batch from our dataloader.
#
# As we unpack the batch, we'll also copy each tensor to the GPU using
# the `to` method.
#
# `batch` contains three pytorch tensors:
#   [0]: input ids
#   [1]: attention masks
#   [2]: labels
b_input_ids = batch[0].to(device)
b_input_mask = batch[1].to(device)
b_labels_actions = batch[2].to(device)
b_labels_attributes = batch[3].to(device)
b_dialog_ids = batch[4].to(device).detach().cpu().numpy()
b_turn_idx = batch[5].to(device).detach().cpu().numpy()

# Tell pytorch not to bother with constructing the compute graph during
# the forward pass, since this is only needed for backprop (training).
with torch.no_grad():
    # Forward pass, calculate logit predictions.
    # token_type_ids is the same as the "segment ids", which
    # differentiates sentence 1 and 2 in 2-sentence tasks.
    result = model(b_input_ids,mask=b_input_mask)

actions_logits_foracc=result['actions'].detach().cpu().numpy()
attributes_logits_foracc=result['attributes'].detach().cpu().numpy()
actions_labels_foracc= b_labels_actions.to('cpu').numpy()
attributes_labels_foracc =b_labels_attributes.to('cpu').numpy()

# Calculate the accuracy for this batch of test sentences, and
# accumulate it over all batches.
accuracy_classification = flat_accuracy_actions(actions_logits_foracc,↵
↵actions_labels_foracc)
accuracy_multilabel = flat_accuracy_attributes(attributes_logits_foracc,↵
↵attributes_labels_foracc)

total_eval_accuracy_classification['matched'] +=↵
↵accuracy_classification['matched']
total_eval_accuracy_classification['counts'] +=↵
↵accuracy_classification['counts']
total_eval_accuracy_multilabel['matched'] += accuracy_multilabel['matched']
total_eval_accuracy_multilabel['counts'] += accuracy_multilabel['counts']

# Fill dictionary for action_evaluation

```

```

for el_i in range(len(actions_logits_foracc)):
    dialog_id = b_dialog_ids[el_i]
    action_log_prob = {}
    for act_i in range(len(actions_logits_foracc[el_i])):
        #todo: controllare che la probabilità predetta sia in scala logaritmica
        → (?? potrebbe essere fonte di errori)
        action_log_prob[le.classes_[act_i]] = np.
        → log(actions_logits_foracc[el_i][act_i])
        #attributes = {}
        attributes = []
        #attributes_list = np.rint(attributes_logits_foracc[el_i])
        attributes_list = np.array(attributes_logits_foracc[el_i])
        for attr in range(len(attributes_list)):
            attribute = mlb.classes_[attr]
            #attributes[mlb.classes_[attr]] = attributes_list[attr]
            if attributes_list[attr] >= 0.5:
                attributes.append(attribute)
        prediction = {
            'action': le.classes_[np.argmax(actions_logits_foracc[el_i])],
            'action_log_prob': action_log_prob,
            'attributes': {'attributes': attributes},
            'turn_id': b_turn_idx[el_i]
        }
    if dialog_id in model_actions:
        model_actions[dialog_id]['predictions'].append(prediction)
    else:
        predictions = list()
        predictions.append(prediction)
        model_actions[dialog_id] = {
            'dialog_id': dialog_id,
            'predictions': predictions
        }

# Report the final accuracy for this validation

#avg_val_accuracy_classification = total_eval_accuracy_classification /
→ len(validation_data_loader)
#avg_val_accuracy_multilabel = total_eval_accuracy_multilabel /
→ len(validation_data_loader)
avg_val_accuracy_classification = total_eval_accuracy_classification['matched']
→ / total_eval_accuracy_classification['counts']
avg_val_accuracy_multilabel = total_eval_accuracy_multilabel['matched']
→ total_eval_accuracy_multilabel['counts']
print(" Accuracy for classification (actions): {0:.4f}".
      → format(avg_val_accuracy_classification))

```

```

print("  Accuracy for multilabel-classification (attributes): {0:.4f}".
      ↪format(avg_val_accuracy_multilabel))

# Reference implementation: evaluation of action prediction along with
↪attributes
metrics = evaluation.evaluate_action_prediction(devtest_dials, model_actions.
      ↪values())
# print("model_actions passed to the evaluator:")
# for v in model_actions.values():
#     print(v)
print("*****")
print("Reference evaluation metrics:")
print(metrics)

```

```

Accuracy for classification (actions): 0.8425
Accuracy for multilabel-classification (attributes): 0.9200
#Instances evaluated API: 5397
*****
Reference evaluation metrics:
{'action_accuracy': 0.8425050954233834, 'action_perplexity': 3.1634650211610884,
'attribute_accuracy': 0.7370522617939476, 'confusion_matrix': array([[ 747.,
53.,      8.,      4.,     13.],
[ 26., 1058.,     91.,    15.,    32.],
[ 12., 179.,  727.,    88.,    28.],
[  0.,   0.,   0.,   0.,   0.],
[  8.,   96.,  118.,   79., 2015.]])}

```

```
[44]: import pandas as pd
```

```
[45]: # Convert test data to dataframe
df_test = pd.DataFrame(data = test_batch)
df_test.head()
```

```
[45]:
```

	ephoc	batchnum	actions_logits \
0	1	1	[[0.0004531708, 0.99980944, 0.0025453316, 0.00...
1	1	2	[[0.00021071802, 0.9996798, 0.0008664634, 0.00...
2	1	3	[[0.0007729576, 0.0009535144, 0.0010184883, 0...
3	1	4	[[0.99967086, 0.0019957258, 0.0006378222, 0.00...
4	1	5	[[7.279184e-05, 0.00031680457, 0.00024304114, ...

	actions_labels \
0	[1, 4, 2, 4, 0, 1, 1, 4, 4, 0, 1, 2]
1	[1, 4, 2, 2, 0, 1, 2, 1, 1, 1, 1, 4]
2	[4, 1, 1, 1, 4, 2, 1, 2, 0, 1, 1, 1]
3	[0, 0, 1, 2, 0, 1, 4, 4, 0, 1, 1, 4]
4	[1, 4, 2, 4, 0, 3, 1, 1, 2, 4, 1, 1]

```

                                attributes_logits \
0 [[0.00041479623, 0.0002373658, 0.001900391, 0...
1 [[0.00058873027, 0.00036660058, 0.0032797686, ...
2 [[0.00016869778, 0.00028577558, 0.90405023, 0...
3 [[0.0009995383, 0.0007337396, 0.0064695496, 0...
4 [[0.00013425016, 0.00026139984, 0.040037967, 0...

                                attributes_labels \
0 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
1 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
2 [[0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
3 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
4 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...

    accuracy_classification      accuracy_multilabel
0 {'matched': 10, 'counts': 12} {'matched': 11, 'counts': 12}
1 {'matched': 9, 'counts': 12}  {'matched': 11, 'counts': 12}
2 {'matched': 8, 'counts': 12}  {'matched': 11, 'counts': 12}
3 {'matched': 9, 'counts': 12}  {'matched': 11, 'counts': 12}
4 {'matched': 9, 'counts': 12}  {'matched': 11, 'counts': 12}

```

```

[46]: # Display floats with two decimal places.
pd.set_option('precision', 2)

# Create a DataFrame from our training statistics.
df_stats = pd.DataFrame(data=training_stats)

# Use the 'epoch' as the row index.
df_stats = df_stats.set_index('epoch')

# A hack to force the column headers to wrap.
#df = df.style.set_table_styles([dict(selector="th",props=[('max-width',
↪ '70px'))])]

# Display the table.
df_stats

```

```

[46]:      Training Loss  Valid. Loss  Valid. Accur. class. \
epoch
1           1.13         1.07         0.85
2           1.06         1.06         0.85
3           1.04         1.05         0.85
4           1.03         1.05         0.86

      Valid. Accur. mult.label Training Time Validation Time
epoch
1           0.88         0:08:38         0:00:28

```

2	0.92	0:08:38	0:00:27
3	0.92	0:08:37	0:00:27
4	0.93	0:08:39	0:00:27

```
[47]: import time

# Objects serialization
timestr = time.strftime("%Y%m%d-%H%M%S")
testdata_filename = f"testdata-{timestr}"
stats_filename = f"stats-{timestr}"
#outtest = open(testdata_filename, "wb")
#outstats = open(stats_filename, "wb")
#pk.dump(obj=df_test, file=outtest)
#outtest.close()
#pk.dump(obj=df_stats, file=outstats)
#outstats.close()
df_test.to_pickle(testdata_filename)
df_stats.to_pickle(stats_filename)
```

```
[48]: import pandas as pd

# Test reimport data
df_stats_reload = pd.read_pickle(stats_filename)
df_test_reload = pd.read_pickle(testdata_filename)

print(testdata_filename)
print(df_stats_reload.head())
print(df_test_reload.head())
```

testdata-20210711-100903

	Training Loss	Valid. Loss	Valid. Accur. class.	\
epoch				
1	1.13	1.07	0.85	
2	1.06	1.06	0.85	
3	1.04	1.05	0.85	
4	1.03	1.05	0.86	

	Valid. Accur. mult.label	Training Time	Validation Time
epoch			
1	0.88	0:08:38	0:00:28
2	0.92	0:08:38	0:00:27
3	0.92	0:08:37	0:00:27
4	0.93	0:08:39	0:00:27

epoch	batchnum	actions_logits \
0	1	1 [[0.0004531708, 0.99980944, 0.0025453316, 0.00...
1	1	2 [[0.00021071802, 0.9996798, 0.0008664634, 0.00...
2	1	3 [[0.0007729576, 0.0009535144, 0.0010184883, 0...

```

3      1      4 [[0.99967086, 0.0019957258, 0.0006378222, 0.00...
4      1      5 [[7.279184e-05, 0.00031680457, 0.00024304114, ...

```

```

                actions_labels \
0 [1, 4, 2, 4, 0, 1, 1, 4, 4, 0, 1, 2]
1 [1, 4, 2, 2, 0, 1, 2, 1, 1, 1, 1, 4]
2 [4, 1, 1, 1, 4, 2, 1, 2, 0, 1, 1, 1]
3 [0, 0, 1, 2, 0, 1, 4, 4, 0, 1, 1, 4]
4 [1, 4, 2, 4, 0, 3, 1, 1, 2, 4, 1, 1]

```

```

                attributes_logits \
0 [[0.00041479623, 0.0002373658, 0.001900391, 0...
1 [[0.00058873027, 0.00036660058, 0.0032797686, ...
2 [[0.00016869778, 0.00028577558, 0.90405023, 0...
3 [[0.0009995383, 0.0007337396, 0.0064695496, 0...
4 [[0.00013425016, 0.00026139984, 0.040037967, 0...

```

```

                attributes_labels \
0 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
1 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
2 [[0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
3 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
4 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...

```

```

                accuracy_classification      accuracy_multilabel
0 {'matched': 10, 'counts': 12} {'matched': 11, 'counts': 12}
1 {'matched': 9, 'counts': 12} {'matched': 11, 'counts': 12}
2 {'matched': 8, 'counts': 12} {'matched': 11, 'counts': 12}
3 {'matched': 9, 'counts': 12} {'matched': 11, 'counts': 12}
4 {'matched': 9, 'counts': 12} {'matched': 11, 'counts': 12}

```

```

[52]: import matplotlib.pyplot as plt
      %% matplotlib inline

      import seaborn as sns

      # Use plot styling from seaborn.
      sns.set(style='darkgrid')

      # Increase the plot size and font size.
      sns.set(font_scale=1.5)
      plt.rcParams["figure.figsize"] = (12,6)

      # Plot the learning curve.
      plt.plot(df_stats['Training Loss'], 'b-o', label="Training")
      plt.plot(df_stats['Valid. Loss'], 'g-o', label="Validation")

```

```
# Label the plot.  
plt.title("Training & Validation Loss")  
plt.xlabel("Epoch")  
plt.ylabel("Loss")  
plt.legend()  
plt.xticks([1, 2, 3, 4])  
  
plt.show()
```



[]: